

Integrating NVDLA into TVM

CS6280 Project Presentation




Dan Wu | Shivam Aggarwal


Nov 19, 2020

Outline

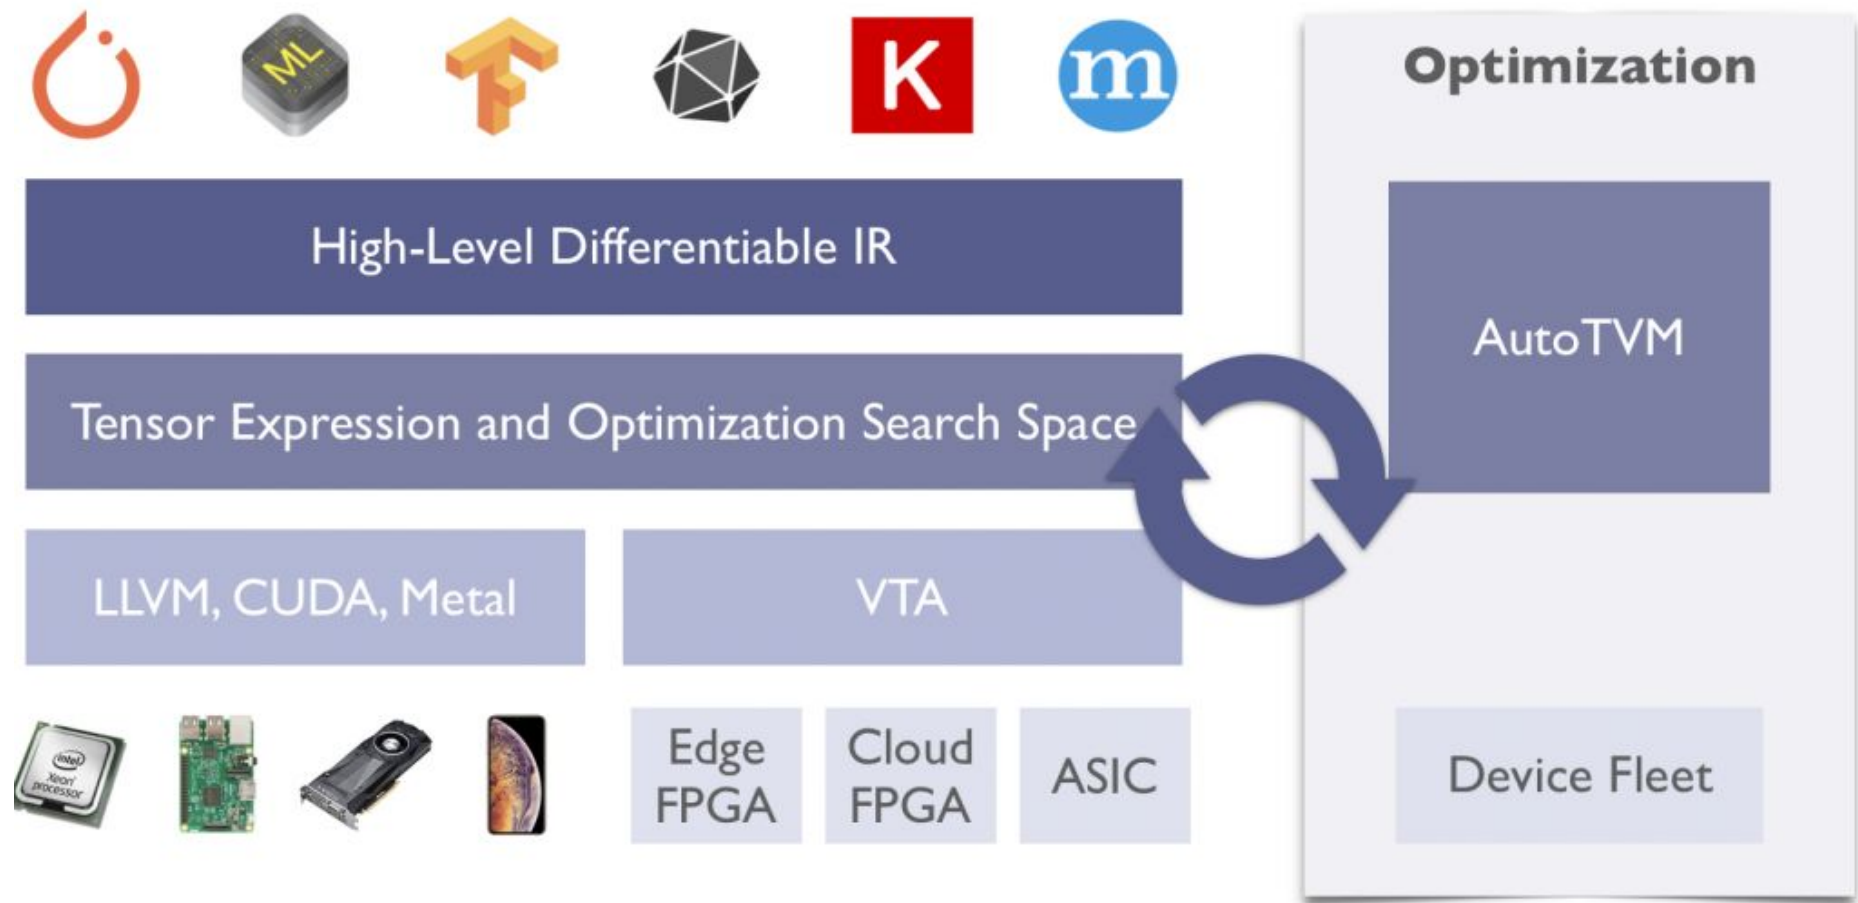
- Background
- Motivation
- Related work
- System design
- Challenges
- Limitations
- Evaluation
- Future work



 **tvm**

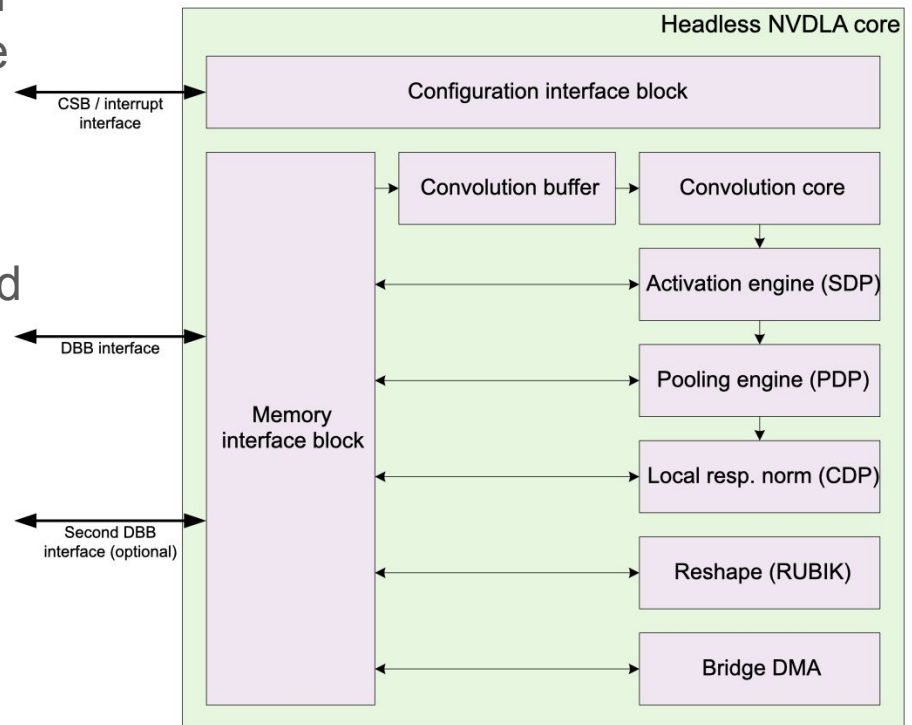


TVM: High level overview

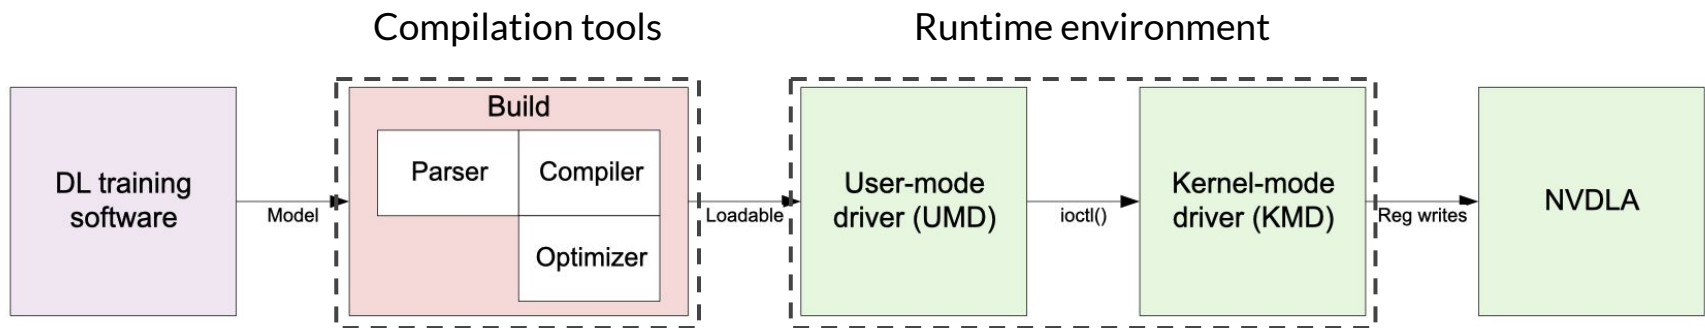


NVDLA: Hardware Architecture

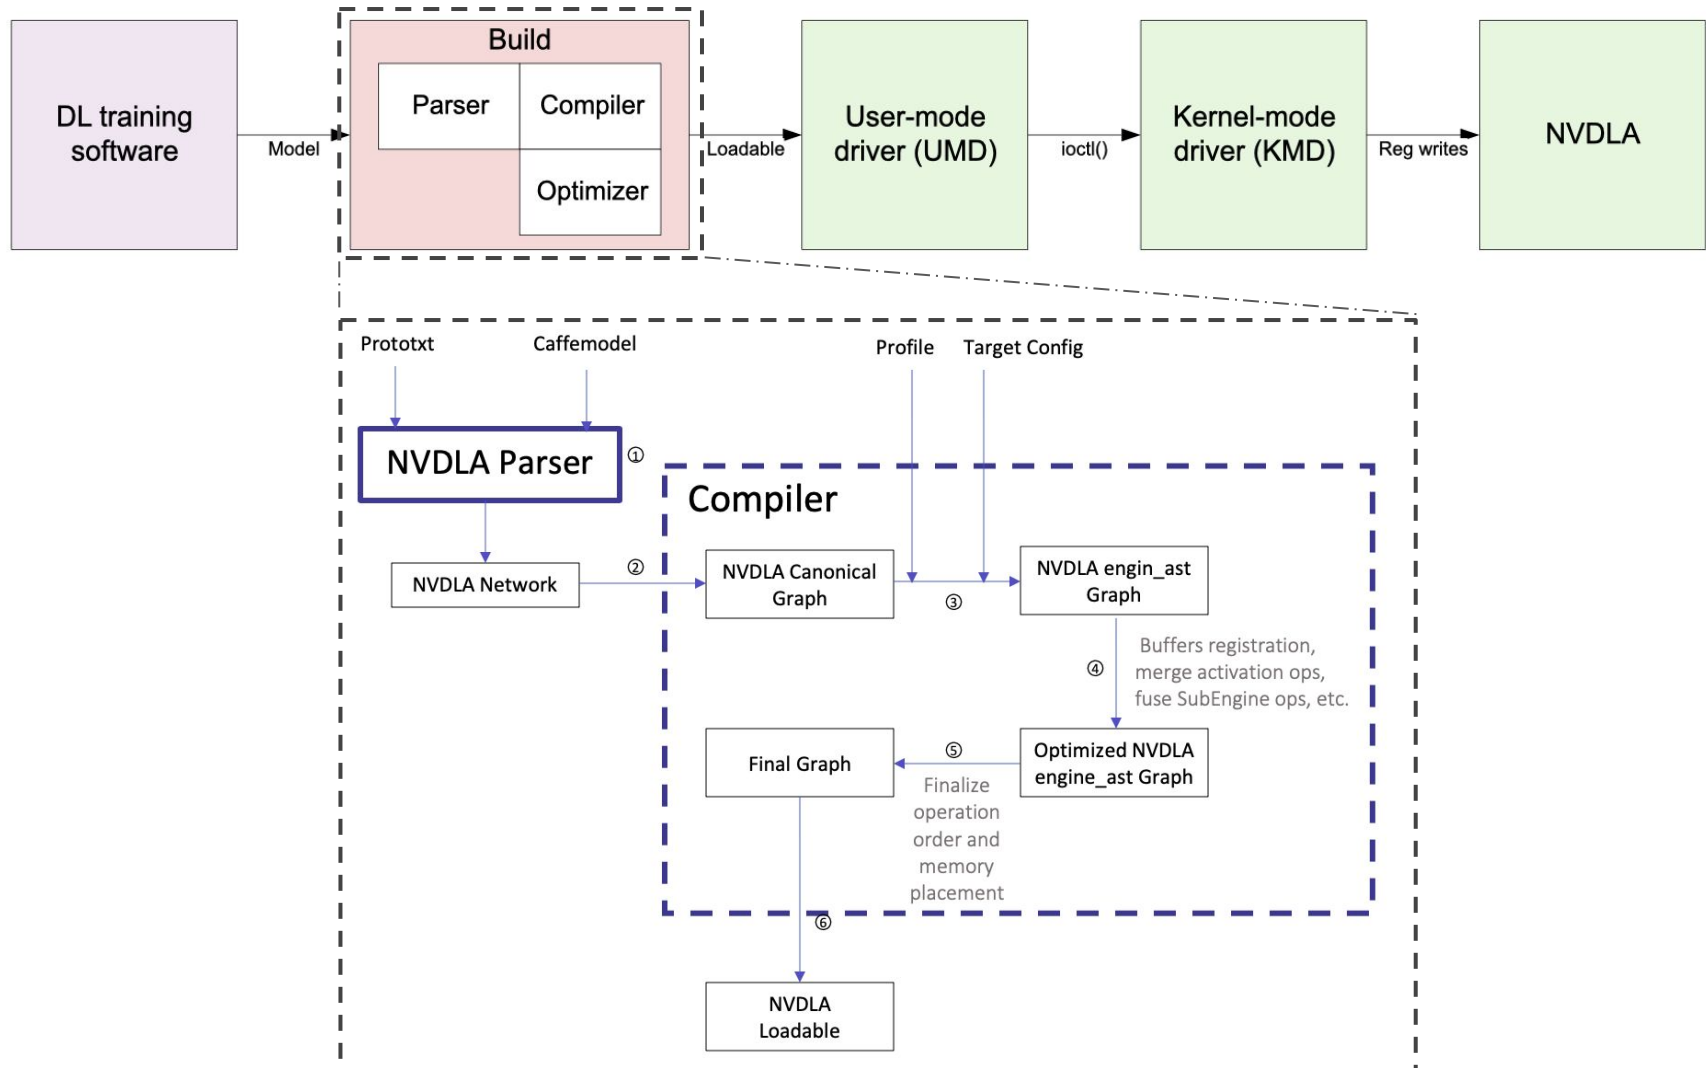
- **Open-source DNN accelerator:** can achieve industry-level DNN inference throughput.
- **Configurable:** hardware configuration, the system's CPU and memory controller configurations, and the application's custom neural network.
- **Scalable:** nv_small, nv_large
- Convolutional core: matrix-matrix multiplication
- Post-processing: activation function, pooling, etc.



NVDLA Workflow



How NVDLA compiler works?



Motivation

NVDLA

Limited Frontend Support

Limited Optimization

No fallback system

NVDLA.org

TVM

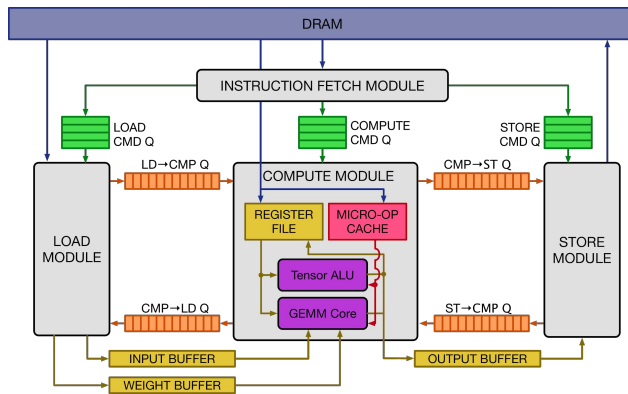
Diverse Frontend Support

Diverse Optimization

Annotation in BYOC



Related work



NVDLA

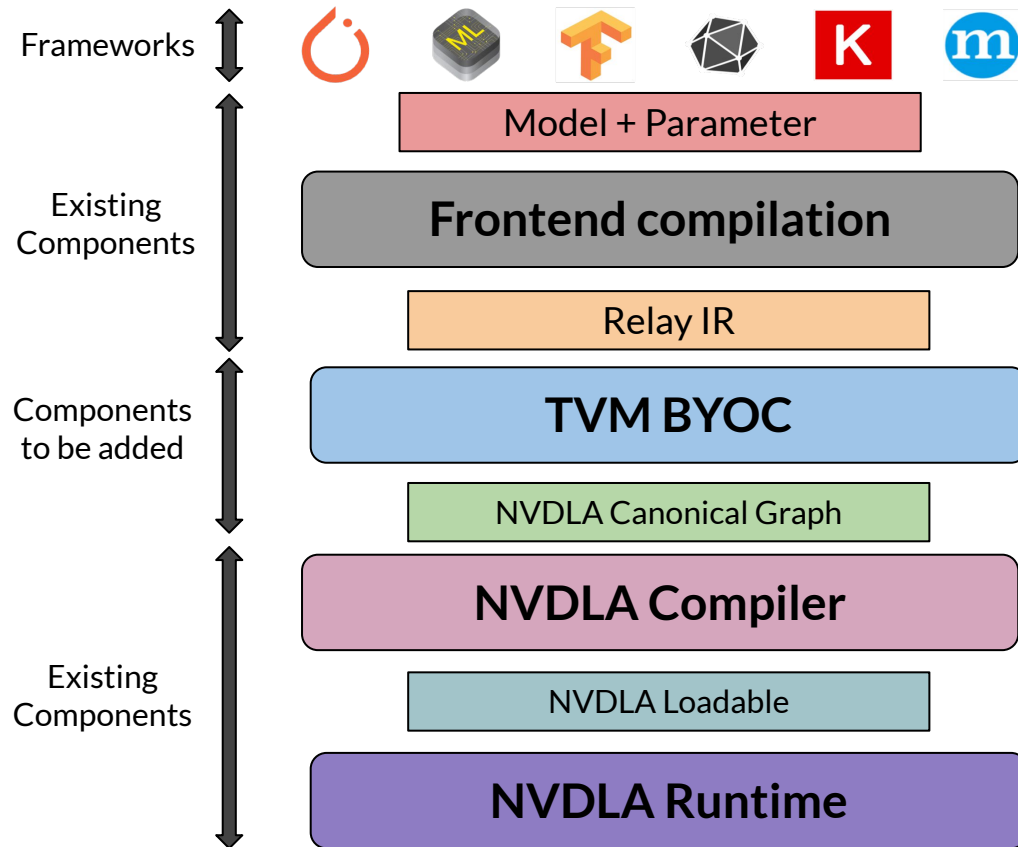


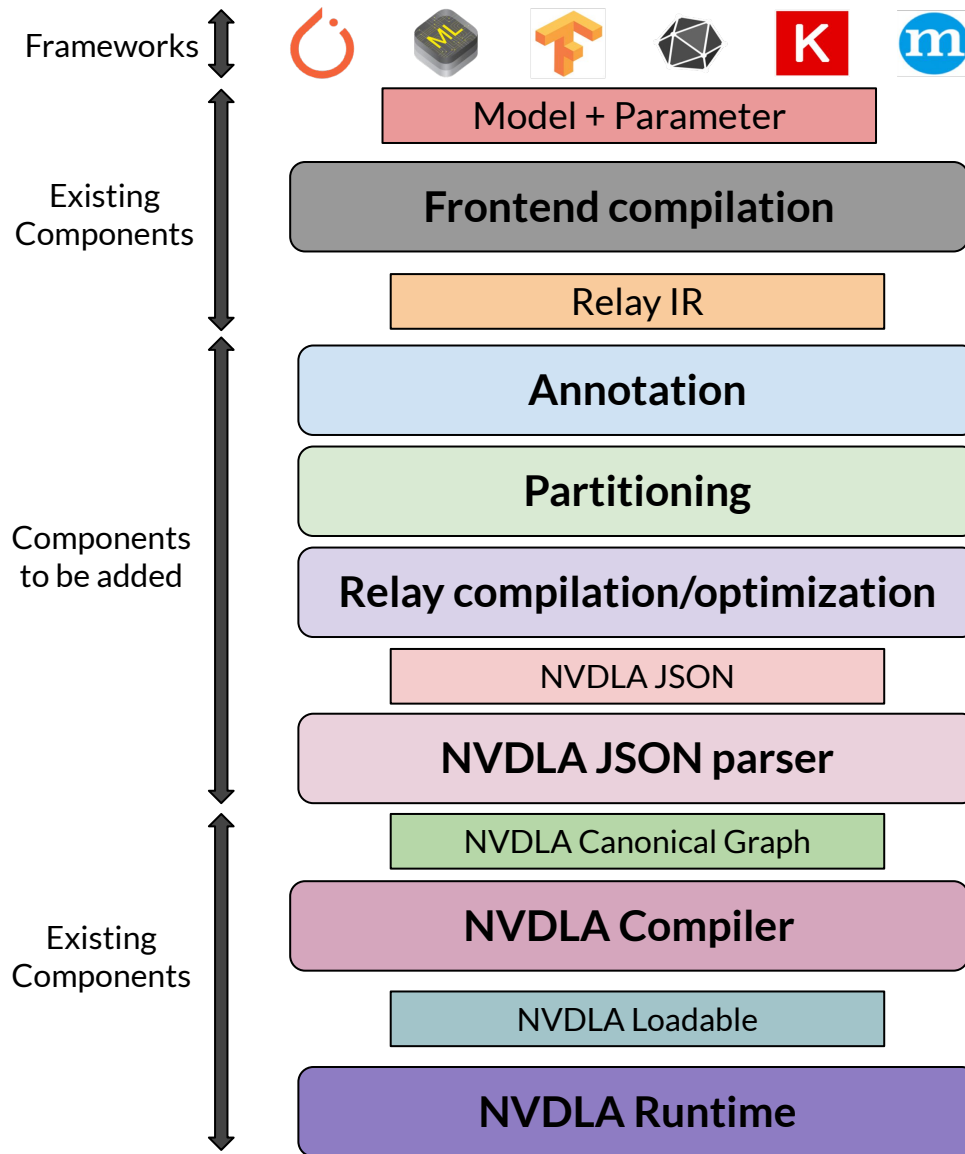
open source DLA

open source compiler



How to integrate the two?





Challenges

- Diverse workflow design
- Different network operator representation
 - In Relay, bias and inputs are stored as network layers;
 - In NVDLA, bias is a parameter of layer and inputs are represented as edges.
- Strict Check in NVDLA compiler
 - Surface check disables single layer testing;
 - Quantization check disallows simple layer without bias item.
- Non-Debuggable NVDLA runtime
 - Older version is more stable but not open-source;
 - Latest version is open-source but have fatal error.

Limitations

- Only high-level optimizations in TVM stack are used.
 - How to involve more lower-level optimizations in TVM?
- No good solution for evaluation.
 - Is it possible to compare the loadable files?
- Not end-to-end process.
 - How to integrate the two code base?

Evaluation

- We can generate JSON file for any trained neural network on diverse ML frameworks using TVM BYOC.
- We can parse a JSON file, rebuild the canonical graph and generate loadable file for all hardware configurations (4 types) in NVDLA.
- We do not have a real NVDLA board and do tests on SystemC-based or AWS FPGA-based simulators.
- Correct output for single convolutional layer.

Future work

- Compare loadable for functional correctness evaluation.
- End-to-end compilation framework for NVDLA.
- Integrating TVM with CGRA based accelerators such as HyCUBE

References

- Chen, Tianqi, Thierry Moreau, Ziheng Jiang, Lianmin Zheng, Eddie Yan, Meghan Cowan, Haichen Shen, et al. “TVM: An Automated End-to-End Optimizing Compiler for Deep Learning.” *ArXiv:1802.04799 [Cs]*, October 5, 2018. <http://arxiv.org/abs/1802.04799>.
- Chen, Tianqi, Lianmin Zheng, Eddie Yan, Ziheng Jiang, Thierry Moreau, Luis Ceze, Carlos Guestrin, and Arvind Krishnamurthy. “Learning to Optimize Tensor Programs.” *ArXiv:1805.08166 [Cs, Stat]*, January 8, 2019. <http://arxiv.org/abs/1805.08166>.
- Moreau, Thierry, Tianqi Chen, Luis Vega, Jared Roesch, Eddie Yan, Lianmin Zheng, Josh Fromm, et al. “A Hardware-Software Blueprint for Flexible Deep Learning Specialization.” *ArXiv:1807.04188 [Cs, Stat]*, April 22, 2019. <http://arxiv.org/abs/1807.04188>.
- Ragan-Kelley, Jonathan, Connelly Barnes, and Andrew Adams. “Halide: A Language and Compiler for Optimizing Parallelism, Locality, and Recomputation in Image Processing Pipelines,” n.d., 12.

Thank you!

Appendix: Relay IR to JSON

```
def @main(%data: Tensor[(1, 14, 14, 512), float32]) -> Tensor[(1, 7, 7, 512), float32] {
  nn.max_pool2d(%data, pool_size=[2, 2], strides=[2, 2], padding=[0, 0, 0, 0],
  layout="NHWC") /* ty=Tensor[(1, 7, 7, 512), float32] */
}
```

```
{
  "input": {
    "dtype": [
      [
        "float32"
      ]
    ],
    "shape": [
      [
        [
          1,
          14,
          14,
          512
        ]
      ]
    ]
  },
  "max_pool2d_0": {
    "ceil_mode": [
      [
        "0"
      ]
    ],
    "dtype": [
      [
        "float32"
      ]
    ],
    "layout": [
      [
        "NHWC"
      ]
    ],
    "num_inputs": "1",
    "num_outputs": "1",
    "padding": [
      [
        "0",
        "0",
        "0",
        "0"
      ]
    ]
  },
  "pool_size": [
    [
      "2",
      "2"
    ]
  ],
  "shape": [
    [
      1,
      7,
      7,
      512
    ]
  ],
  "strides": [
    [
      "2",
      "2"
    ]
  ]
}
```